

4. MODEL TRAINING

January 2, 2026

Importing Libraries

```
[2]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, StratifiedKFold

from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from hyperopt import hp, tpe, STATUS_OK, Trials, fmin

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, precision_recall_curve, ConfusionMatrixDisplay

import matplotlib.pyplot as plt
import shap

import pickle

import warnings
warnings.filterwarnings('ignore')
```

Reading the data

```
[3]: df = pd.read_csv('data/processed/processed_data.csv')

df.head()
```

```
[3]:   has_https  url_len  domain_len  path_len  query_len  url_depth \
0           1      25          10         6          0           1
1           1      37          22         6          0           1
2           1      32          18         5          0           1
3           1      46          14         0         22           1
4           1      32          18         5          0           1

  subdomain_count  tld_len  url_has_ipv4  url_has_port ...  spl_char_count \
0                  1        2                 0                 0 ...                 5
```

```

1          1      3      0      0 ...      5
2          1      3      0      0 ...      6
3          1      3      0      0 ...      7
4          1      3      0      0 ...      6

    url_entropy  domain_entropy  sld_entropy  path_entropy  domain_token_count \
0      3.863465      2.721928      2.235926      2.521641              2
1      4.208925      3.629220      3.419382      2.521641              2
2      4.452820      3.947703      3.000000      2.584963              3
3      4.760096      3.664498      3.121928     -0.000000              2
4      4.241729      3.836592      3.000000      2.584963              3

    path_token_count  total_tokens  avg_token_length  class
0                  1            3             5.00       1
1                  1            3             9.00       1
2                  1            4             5.25       1
3                  1            3             8.50       1
4                  1            4             5.25       1

[5 rows x 30 columns]

```

Splitting the data

```
[4]: X = df.drop(columns=['class']).values
y = df['class'].values
```

```
[5]: X_train,X_test,y_train,y_test = train_test_split(X,y,train_size=0.
    ↪75,random_state=6)
```

Training the Models

```
[6]: # Function to evaluate the models
def evaluate_model(y_true,y_pred):
    accuracy = accuracy_score(y_true,y_pred)
    precision = precision_score(y_true,y_pred)
    recall = recall_score(y_true,y_pred)
    f1 = f1_score(y_true,y_pred)

    print(f'Accuracy: {accuracy}')
    print(f'Precision: {precision}')
    print(f'Recall: {recall}')
    print(f'F1-Score: {f1}')
```

Random Forest Classifier

```
[7]: rf_model = RandomForestClassifier()
rf_model.fit(X_train,y_train)

rf_pred_train = rf_model.predict(X_train)
```

```

rf_pred_test = rf_model.predict(X_test)

print('Metrics of Random Forest on Training data:')
evaluate_model(y_train,rf_pred_train)

print()

print('Metrics of Random Forest on Testing data:')
evaluate_model(y_test,rf_pred_test)

```

Metrics of Random Forest on Training data:

Accuracy: 0.9980076354092102
 Precision: 0.99923558711841
 Recall: 0.9966551071103569
 F1-Score: 0.997943678967653

Metrics of Random Forest on Testing data:

Accuracy: 0.9792412312097352
 Precision: 0.9878127247265432
 Recall: 0.9688903738352452
 F1-Score: 0.9782600547246898

XGBoost Classifier

```

[8]: xg_model = XGBClassifier()
xg_model.fit(X_train,y_train)

xg_pred_train = xg_model.predict(X_train)
xg_pred_test = xg_model.predict(X_test)

print('Metrics of XGBoost Classifier on Training data:')
evaluate_model(y_train,xg_pred_train)

print()

print('Metrics of XGBoost Classifier on Testing data:')
evaluate_model(y_test,xg_pred_test)

```

Metrics of XGBoost Classifier on Training data:

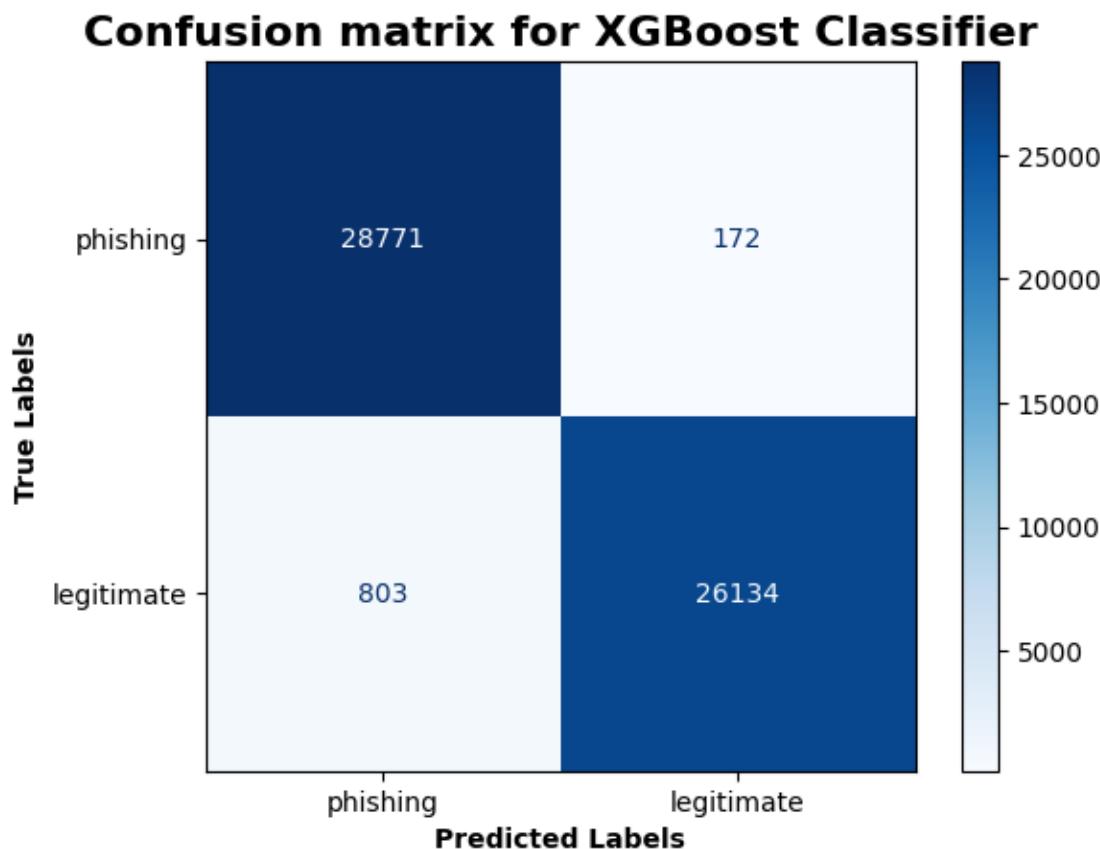
Accuracy: 0.9858088761632069
 Precision: 0.9963655570506935
 Recall: 0.974298433311198
 F1-Score: 0.9852084434358194

Metrics of XGBoost Classifier on Testing data:

Accuracy: 0.9825518969219756
 Precision: 0.9934615677031856
 Recall: 0.9701897018970189
 F1-Score: 0.9816877335987829

By comparing the metrics of Random Forest Classifier & XGBoost Classifier on both training set and testing set, XGBoost Classifier is best. Both the model's metrics are good on training set but on testing set, Random Forest model's metrics are slightly lower than XGBoost model's metrics telling that Random Forest model has less Generalization. Therefore, we prefer XGBoost Classifier.

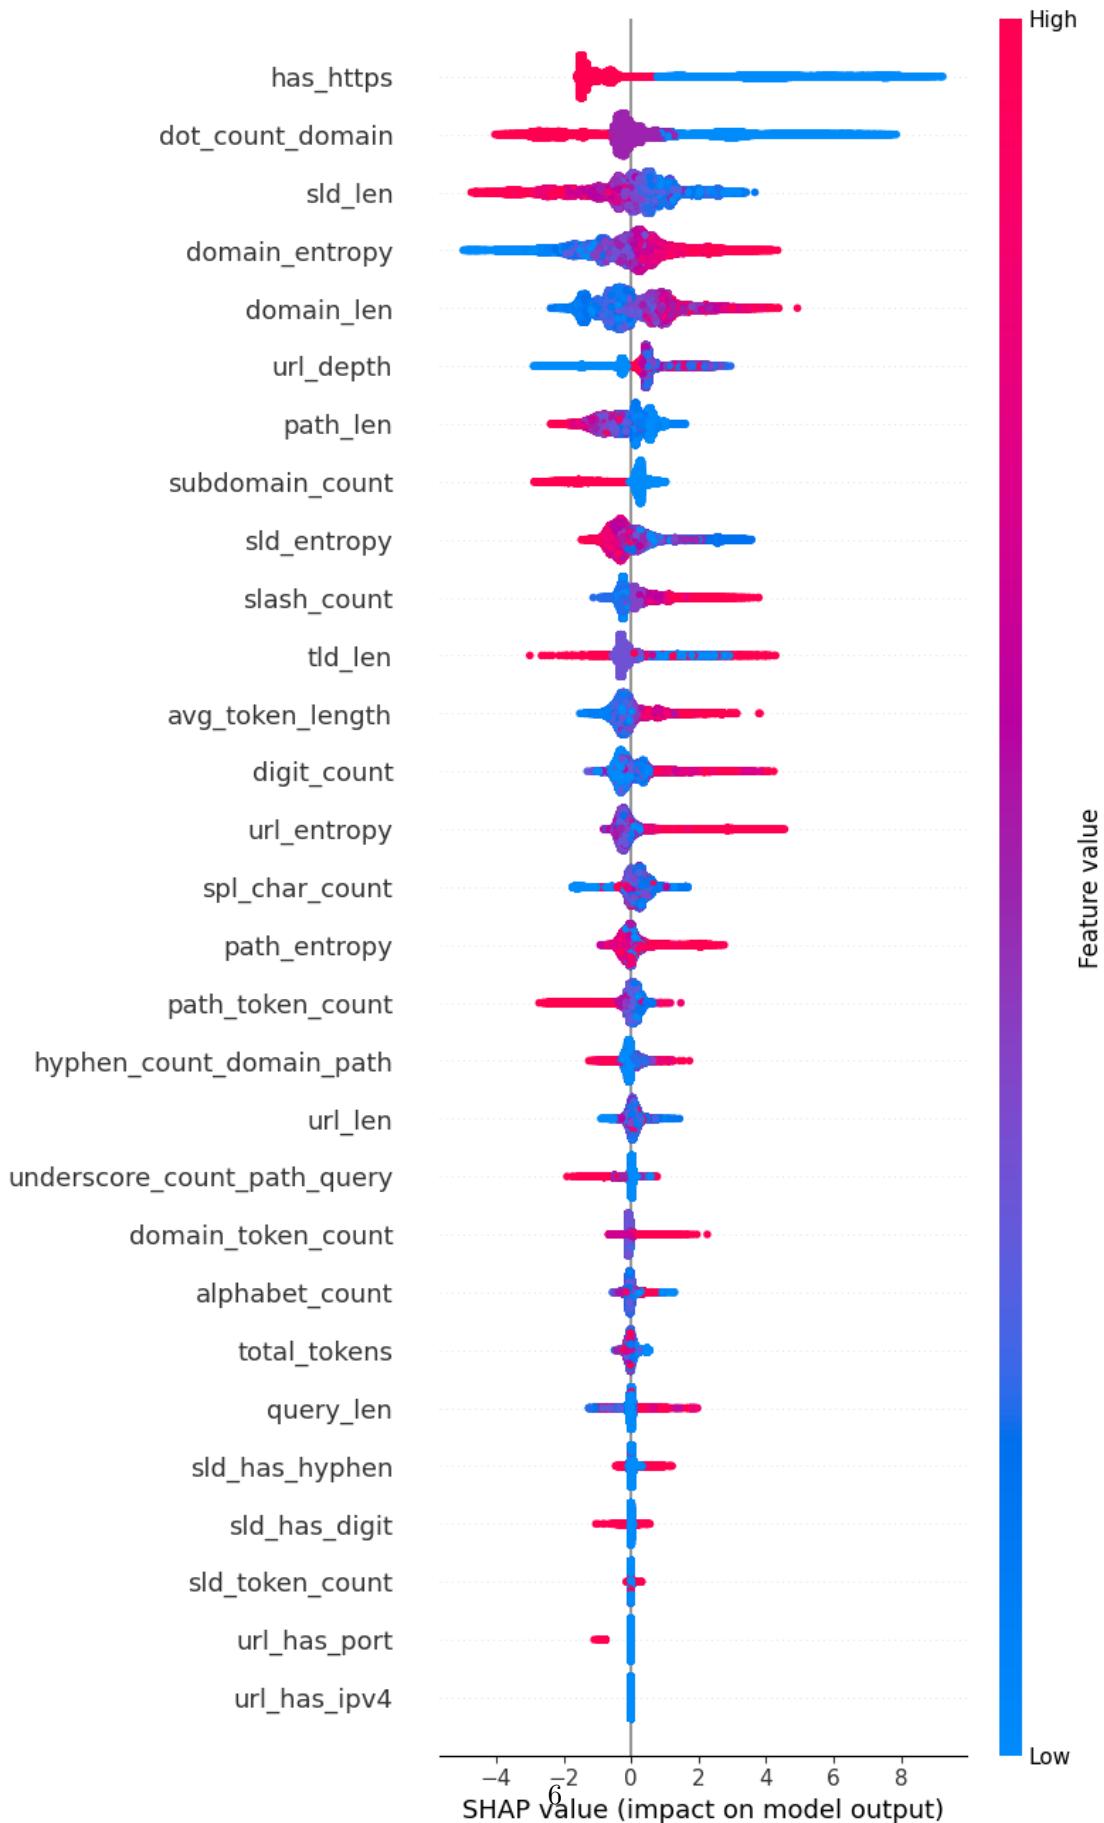
```
[9]: disp = ConfusionMatrixDisplay.  
      ↪from_estimator(xg_model,X_test,y_test,display_labels=['phishing','legitimate'],cmap=plt.  
      ↪cm.Blues)  
plt.xlabel('Predicted Labels',weight='bold')  
plt.ylabel('True Labels',weight='bold')  
plt.title('Confusion matrix for XGBoost Classifier',weight='bold',fontsize=16)  
plt.show()
```



Feature Importance

```
[10]: features = df.drop(columns='class').columns  
explainer = shap.TreeExplainer(xg_model)  
shap_values = explainer.shap_values(X_test)  
  
shap.initjs()
```

```
<IPython.core.display.HTML object>  
[11]: X_test_df = pd.DataFrame(X_test,columns=features)  
shap.summary_plot(shap_values,X_test_df,max_display=30)
```



From the above plot, most of the features have some impact in predicting phishing URLs. But there are some features which does not show any impact on predictions: - total_tokens - sld_token_count - url_has_ipv4 - url_has_port

So, we will ignore these features.

Therefore the final features considered to train the model are: - has_https - url_len - domain_len - path_len - query_len - url_depth - subdomain_count - tld_len - sld_len - sld_has_digit - sld_has_hyphen - dot_count_domain - hyphen_count_domain_path - underscore_count_path_query - slash_count - digit_count - alphabet_count - spl_char_count - url_entropy - domain_entropy - sld_entropy - path_entropy - domain_token_count - path_token_count - avg_token_length

```
[12]: # Updated data
updated_df = df.
    ↪drop(columns=['total_tokens','sld_token_count','url_has_ipv4','url_has_port'])

X_updated = updated_df.drop(columns='class').values
y_updated = updated_df['class']

X_train_updated,X_test_updated,y_train_updated,y_test_updated =
    ↪train_test_split(X_updated,y_updated,train_size=0.75,random_state=6)
```

```
[13]: # Saving the updated data
updated_df.to_csv('data/feature_refined/feature_refined_data.csv',index=False)
print('Updated data saved')
```

Updated data saved

```
[14]: model = XGBClassifier()
model.fit(X_train_updated,y_train_updated)
```

```
[14]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[15]: print('Metrics of XGBoost Classifier on Training data:')
evaluate_model(y_train_updated,model.predict(X_train_updated))
```

```

print()

print('Metrics of XGBoost Classifier on Testing data:')
evaluate_model(y_test_updated,model.predict(X_test_updated))

```

Metrics of XGBoost Classifier on Training data:

Accuracy: 0.9859520400858983
 Precision: 0.9962169295544523
 Recall: 0.9747411397230625
 F1-Score: 0.9853620332788424

Metrics of XGBoost Classifier on Testing data:

Accuracy: 0.9828203292770222
 Precision: 0.9936528448177568
 Recall: 0.9705609384860971
 F1-Score: 0.9819711538461539

Hyperparameter Tuning

```
[16]: search_space = {
    "max_depth" : hp.quniform("max_depth",3,10,1),
    "learning_rate" : hp.loguniform("learning_rate",np.log(0.01),np.log(0.2)),
    "n_estimators" : hp.quniform("n_estimators",100,600,50),
    "subsample" : hp.uniform("subsample",0.6,1.0),
    "colsample_bytree" : hp.uniform("colsample_bytree",0.6,1.0),
    "gamma" : hp.uniform("gamma",0,5),
    "min_child_weight" : hp.qloguniform("min_child_weight",1,10,1),
    "reg_alpha" : hp.loguniform("reg_alpha",np.log(1e-3),np.log(1)),
    "reg_lambda" : hp.loguniform("reg_lambda",np.log(1),np.log(10))
}
```

```
[17]: def objective(params):

    params['max_depth'] = int(params['max_depth'])
    params['n_estimators'] = int(params['n_estimators'])
    params['min_child_weight'] = int(params['min_child_weight'])

    params.update({
        "objective" : "binary:logistic",
        "eval_metric" : "logloss",
        "random_state" : 42,
        "tree_method" : "hist",
        "n_jobs" : -1
    })

    skf = StratifiedKFold(n_splits=5,shuffle=True,random_state=42)

    recalls = []
```

```

for train_idx, val_idx in skf.split(X_train, y_train):
    X_tr, X_val = X_train[train_idx], X_train[val_idx]
    y_tr, y_val = y_train[train_idx], y_train[val_idx]

    model = XGBClassifier(**params)
    model.fit(X_tr, y_tr)

    y_pred = model.predict(X_val)
    recalls.append(recall_score(y_val, y_pred))

mean_recall = np.mean(recalls)
std_recall = np.std(recalls)

return {"loss": -mean_recall, "status": STATUS_OK}

```

[18]: trials = Trials()

```

best_params = fmin(
    fn = objective,
    space = search_space,
    algo = tpe.suggest,
    max_evals = 50,
    trials = trials
)

```

100% | 50/50 [04:07<00:00, 4.96s/trial, best loss:
-0.9702648921764627]

[19]: best_params

```

{'colsample_bytree': 0.6467775791690556,
'gamma': 1.1327649937229542,
'learning_rate': 0.19739238497749395,
'max_depth': 6.0,
'min_child_weight': 3.0,
'n_estimators': 300.0,
'reg_alpha': 0.05223219666480833,
'reg_lambda': 2.2337823495400113,
'subsample': 0.6033832013130304}

```

[20]: best_params['max_depth'] = int(best_params['max_depth'])
best_params['n_estimators'] = int(best_params['n_estimators'])
best_params['min_child_weight'] = int(best_params['min_child_weight'])

best_params.update({
 'objective': 'binary:logistic',
 'eval_metric': 'logloss',
})

```

        'random_state' : 42,
        'n_jobs' : -1
    })

final_model = XGBClassifier(**best_params)
final_model.fit(X_train,y_train)

print('Metrics of XGBoost Classifier on Training data:')
evaluate_model(y_train,final_model.predict(X_train))

print()

print('Metrics of XGBoost Classifier on Testing data:')
evaluate_model(y_test,final_model.predict(X_test))

```

Metrics of XGBoost Classifier on Training data:

Accuracy: 0.9871391076115485
 Precision: 0.9960149377177373
 Recall: 0.9773973781942498
 F1-Score: 0.9866183370987364

Metrics of XGBoost Classifier on Testing data:

Accuracy: 0.9824087329992842
 Precision: 0.9920376128004853
 Recall: 0.9713034116642536
 F1-Score: 0.9815610286807601

Handling False Negatives

```
[30]: import numpy as np
from sklearn.metrics import precision_recall_curve, f1_score

def recall_and_threshold_constrained_selection(
    y_true,
    y_prob,
    min_recall=0.9725,
    min_threshold=0.45
):
    precision, recall, thresholds = precision_recall_curve(y_true, y_prob)

    # precision_recall_curve returns one extra precision/recall value
    precision = precision[:-1]
    recall = recall[:-1]

    thresholds = np.array(thresholds)

    # Apply BOTH constraints
    valid_idxs = np.where(

```

```

        (recall >= min_recall) &
        (thresholds >= min_threshold)
    )[0]

    if len(valid_idxs) == 0:
        raise ValueError(
            f"No threshold satisfies recall >= {min_recall} "
            f"and threshold >= {min_threshold}"
        )

    # Among valid thresholds, maximize precision
    best_idx = valid_idxs[np.argmax(precision[valid_idxs])]
    best_threshold = thresholds[best_idx]

    return {
        "threshold": float(best_threshold),
        "precision": float(precision[best_idx]),
        "recall": float(recall[best_idx]),
        "f1": float(
            f1_score(
                y_true,
                (y_prob >= best_threshold).astype(int)
            )
        ),
        "accuracy": accuracy_score(
            y_true,
            (y_prob >= best_threshold).astype(int)
        )
    }
}

```

[22]: y_prob = final_model.predict_proba(X_test)[:,1]
y_prob

[22]: array([0.9997161 , 0.00365725, 0.02950696, ..., 0.99990857, 0.9999615 ,
0.99667823], shape=(55880,), dtype=float32)

[31]: recall_and_threshold_constrained_selection(y_test,y_prob)

[31]: {'threshold': 0.46429693698883057,
'precision': 0.9904347826086957,
'recall': 0.9725284924082117,
'f1': 0.9813999662839268,
'accuracy': 0.9822297780959198}

After applying a recall-constrained threshold (0.45) slightly increased recall to 0.9725 while maintaining very high precision (0.9904), resulting in an unchanged F1-score (0.9821) and accuracy (0.9822) compared to the initial model. Now the threshold is 0.464.

Saving the model

```
[32]: with open('models/XGBoostModel.pkl','wb') as file:  
    pickle.dump(final_model,file=file)
```